

Privacy Proofs for OpenDP: MakeCount

Connor Wagaman – wagaman@college.harvard.edu

July 16, 2021

Contents

1	MakeCount	1
1.1	Implementation of MakeCount in Rust	1
1.2	Implementation of MakeCount in Python-style pseudocode, with preconditions	2
2	Proofs for the pseudocode	3

1 MakeCount

1.1 Implementation of MakeCount in Rust

In OpenDP (Rust), this is called `make_count`. See <https://github.com/opensdp/opensdp/blob/main/rust/opensdp/src/trans/count.rs>.

This proof is based on the code in <https://github.com/opensdp/opensdp/blob/c3b5c3bd9fc50c556362b628f08c5fddea069b4d/rust/opensdp/src/trans/count.rs#L14-L27> from 12 July 2021. (It is from [this pull request](#).) The Rust code can also be seen below.

```
1 pub fn make_count<TIA, T0>(  
2 ) -> Fallible<Transformation<VectorDomain<AllDomain<TIA>>, AllDomain<T0>, SymmetricDistance, AbsoluteDistance<T0>>>  
3  
4     where T0: ExactIntCast<usize> + One + DistanceConstant<IntDistance>, IntDistance: InfCast<T0> {  
5  
6     Ok(Transformation::new(  
7         VectorDomain::new_all(),  
8         AllDomain::new(),
```

```

9      // think of this as: min(arg.len(), T0::MAX_CONSECUTIVE)
10     Function::new(move |arg: &Vec<TIA>|
11         T0::exact_int_cast(arg.len()).unwrap_or(T0::
MAX_CONSECUTIVE)),
12     SymmetricDistance::default(),
13     AbsoluteDistance::default(),
14     StabilityRelation::new_from_constant(T0::one()))
15 }

```

1.2 Implementation of MakeCount in Python-style pseudocode, with preconditions

We now use Python-style pseudocode to present a representation of the Rust function.

Recall that functions in the pseudocode are defined in the document “[List of definitions used in the pseudocode](#)”.

*The use of **code**-style parameters in the preconditions section below (for example, **input_domain**) means that this information should be passed along to the **Transformation** constructor.*

Here, we use preconditions to check for traits, and to specify the domains and metrics.

Preconditions

- **User-specified types:** The `make_count` function takes two inputs: a generic input type `TIA` for the `Transformation` (meaning that the input vector to `Transformation` is of type `Vec(TIA)`), and a generic output type `T0` for the `Transformation`.
 - `T0` has traits `One`, `ExactIntCast(usize)`, and `DistanceConstant(IntDistance)`
 - `IntDistance` has trait `InfCast(T0)`
 - **Question:** The final bullet point above is not needed in this proof, but it is needed in the code so a hint can be constructed (otherwise a binary search would be needed to construct the hint). Should this precondition be included here or not?

Postconditions: a `Transformation` must be returned (i.e. if a `Transformation` cannot be returned successfully, a runtime error should be returned)

```

1 def MakeCount(TIA, T0):
2
3     input_domain = VectorDomain(AllDomain(TIA))
4     output_domain = AllDomain(T0)

```

```

5   input_metric = SymmetricDistance()
6   output_metric = AbsoluteDistance(T0)
7
8   # give the Transformation the following properties
9   max_value = get_max_consecutive_int(T0)
10  def function(data: Vec<TIA>) -> T0:
11      try:
12          return exact_int_cast(len(data), T0)
13      except FailedCast:
14          return max_value
15  def stability_relation(din:u32, dout:T0) -> bool:
16      return 1 * inf_cast(din,T0) <= dout
17
18  # now, return the Transformation
19  return Transformation(input_domain,output_domain,function,
input_metric,output_metric,stability_relation)

```

2 Proofs for the pseudocode

Theorem 2.1. *For every setting of the input parameters TIA , $T0$ for `MakeCount` such that the given preconditions hold, the `Transformation` returned by `MakeCount` has the following properties:*

1. (Appropriate output domain). *For every vector v in the `input_domain`, `function(v)` is in the `output_domain`.*
2. (Domain-metric compatibility). *The domain `input_domain` matches one of the possible domains listed in the definition of `input_metric`, and likewise `output_domain` matches one of the possible domains listed in the definition of `output_metric`.*
3. (Stability guarantee). *For every input u, v drawn from the `input_domain` and for every pair (d_{in}, d_{out}) , where d_{in} is of type `u32` and d_{out} is of type `T0` (see line 15 of the pseudocode), if u, v are d_{in} -close under the `input_metric` and `stability_relation(din, dout) = True`, then `function(u), function(v)` are d_{out} -close under the `output_metric`.*

Proof. **(Part 1 – appropriate output domain).** In section 1.2, we see that any value of type `T0` is in the `output_domain`, and in line 10 of the Python-style pseudocode, we see that the `function` is always guaranteed to return a value of type `T0`. Therefore, since our output domain is any value of type `T0`, we see that `function` has the appropriate output domain `output_domain`.

Moreover, for some input vector v drawn from `input_domain`, `function` either returns `exact_int_cast(len(data), T0)`, which will be of type `T0` by the definition

of `exact_int_cast`; or, if the casting fails, it returns `get_max_consecutive_int(T0)` which, from our definition of `get_max_consecutive_int`, will be of type `T0`. Therefore, since our output domain is always some value of type `T0`, we see that `function` has the appropriate output domain `output_domain`. \square

Question: Is the second paragraph in the proof above of “Appropriate output domain” necessary?

Proof. (Part 2 – domain-metric compatibility).

The `input_domain` is `VectorDomain(AllDomain(TIA))`. Because our `input_metric` of `SymmetricDistance` is compatible with any domain of the form `VectorDomain(inner_domain)`, and because `VectorDomain(AllDomain(TIA))` is of this form, we see that it is compatible with our `input_metric` of `SymmetricDistance`.

The `output_domain` is `AllDomain(T0)`. Because our `output_metric` of `SymmetricDistance` is compatible with any domain of the form `AllDomain(T)` where `T` has the trait `Sub(Output=T)`, and because `AllDomain(T0)` is of this form and has the necessary trait, we see that it is compatible with our `output_metric` of `AbsoluteDistance`. \square

Proof. (Part 3 – stability relation). We consider two inputs: a vector `u` of elements of type `TIA`; and a vector `v` of elements of type `TIA`. (This `input_domain` is specified in the pseudocode in section 1.2.)

Assume it is the case that `stability_relation(din, dout) = True`. From the stability relation provided on line 16, this means that `inf_cast(din, T0) ≤ dout`. Recall that `inf_cast` will cast `din` to a value at least as large as `din`, so this assumption that `stability_relation` is `True` also means that `din ≤ dout`. Also assume that `v`, `w` are `din`-close under the symmetric distance metric (in accordance with the `input_metric` specified in the preconditions in section 1.2).

We now refer to the definition of symmetric distance provided in the [Proof Definitions](#) document; the definition is copied here for convenience:

Definition 2.1 (Symmetric distance). Let u, v be vectors of elements drawn from domain \mathcal{X} . Define $m_v(\ell)$ as the multiplicity of element ℓ in vector v . For example, if v contains five instances of the number “21”, then $m_v(21) = 5$.

A definition of the symmetric distance between u and v , then, is

$$d_{\text{Sym}}(u, v) = \sum_{z \in \mathcal{X}} |m_u(z) - m_v(z)|.$$

Question: How should I refer readers to a definition located in another document? I know how to use `\label{...}` and `\ref{...}`, but that’s only for referring to definitions, sections, etc. located in the same doc.

Question: As a follow-up to the question above, if there's not a good way to refer to definitions in proofs, should important definitions be copied into the proof doc (as above), or should I remove “in-proof” definitions and rely on readers to track down the right definition in the proof definitions document, which is a document that may be continuously updated for the lifetime of the OpenDP project?

Combining the assumptions that $\text{inf_cast}(d_{\text{in}}, T_0) \leq d_{\text{out}}$ and that \mathbf{v}, \mathbf{w} are d_{in} -close under the symmetric distance metric means that

$$d_{\text{Sym}}(\mathbf{u}, \mathbf{v}) \leq d_{\text{in}} \leq d_{\text{out}}. \quad (1)$$

Let \mathcal{X} be the domain of all elements of type **TIA**. Therefore, we see that the symmetric distance between \mathbf{u} and \mathbf{v} is

$$d_{\text{Sym}}(\mathbf{u}, \mathbf{v}) = \sum_{z \in \mathcal{X}} |m_{\mathbf{u}}(z) - m_{\mathbf{v}}(z)| \leq d_{\text{in}} \leq d_{\text{out}}. \quad (2)$$

The **function** used in **MakeCount** sums over a single data type, namely a row. Let **rows** be a one-element domain, where every element of type **TIA** is considered to be the same element of **rows**; and let the single element be called **row**. Also, as in the **Pseudocode definitions** document, let $\text{len}(\text{vec})$ be a function that returns the number of rows in vector **vec**.

Therefore, using the notation in definition 2.1, we can write

$$|\text{len}(\mathbf{u}) - \text{len}(\mathbf{v})| = \sum_{z \in \text{rows}} |m_{\mathbf{u}}(z) - m_{\mathbf{v}}(z)| = |m_{\mathbf{u}}(\text{row}) - m_{\mathbf{v}}(\text{row})| \quad (3)$$

(note that the summation term is removed in the final term in equation 3 since the domain **rows** consists of the single element **row**).

By the triangle inequality, then, we see that

$$|m_{\mathbf{u}}(\text{row}) - m_{\mathbf{v}}(\text{row})| \leq \sum_{z \in \mathcal{X}} |m_{\mathbf{u}}(z) - m_{\mathbf{v}}(z)|. \quad (4)$$

Combining equations 3 and 4 tells us that $|\text{len}(\mathbf{u}) - \text{len}(\mathbf{v})| = |m_{\mathbf{u}}(\text{row}) - m_{\mathbf{v}}(\text{row})| \leq \sum_{z \in \mathcal{X}} |m_{\mathbf{u}}(z) - m_{\mathbf{v}}(z)|$; combining this with equation 2 tells us that we have

$$|\text{len}(\mathbf{u}) - \text{len}(\mathbf{v})| = |m_{\mathbf{u}}(\text{row}) - m_{\mathbf{v}}(\text{row})| \leq d_{\text{out}}, \quad (5)$$

so $\text{len}(\mathbf{u})$ and $\text{len}(\mathbf{v})$ must be d_{out} -close. This, however, does not complete the proof because **function**(\mathbf{u}) does not return $\text{len}(\mathbf{u})$, but either **exact_cast**($\text{len}(\mathbf{u}), T_0$) or – in the event **exact_cast** fails – **get_max_consecutive_int**(T_0).

We now consider the two cases that could occur:

1. (Without loss of generality, `exact_cast(len(u), T0)` fails and `exact_cast(len(v), T0)` succeeds). Because `T0` has trait `ExactIntCast(usize)`, if the `exact_cast` fails for `len(u)`, we then know that `len(u)` is greater than `get_max_consecutive_int(T0)`. Likewise, if the `exact_cast` succeeds for `len(v)`, we then know that `len(v)` is no larger than `get_max_consecutive_int(T0)`. Therefore, because the return value `get_max_consecutive_int(T0)` for `u` is smaller than the true length value `len(u)`, the absolute difference between the output for `u` and the output for `v` will be *smaller* than the absolute distance between `len(u)` and `len(v)`. Since we showed that the `len(u)` and `len(v)` are d_{out} -close in equation 5, therefore the outputs will still be d_{out} -close.

Note that if `exact_cast` fails for both `len(u)` and `len(v)`, then the output for both `u` and `v` is `get_max_consecutive_int(T0)`, resulting in an absolute distance of 0 between the outputs – the smallest possible absolute distance – so the outputs for `u` and `v` must be d_{out} -close.

2. (Both `exact_cast(len(u), T0)` and `exact_cast(len(v), T0)` succeed). Because `T0` implements `ExactIntCast(usize)`, we know `exact_casts` from `len(u)` to `T0` will be exact. Therefore, the returned values will be `len(u)` and `len(v)`, except the values will now be of type `T0`. Since we showed that the `len(u)` and `len(v)` are d_{out} -close in equation 5, therefore the `exact_casted` lengths will also be d_{out} -close.

Because the outputs will always be d_{out} -close for inputs that follow the conditions specified in part 2 of theorem 2.1, we see that the stability guarantee is proven.

□